

```
In [1]: import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt
import random
```

```
In [2]: mnist = tf.keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
In [3]: x_train = x_train/255
x_test = x_test/255
```

```
In [4]: model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(128, activation="relu"),
    keras.layers.Dense(10, activation="softmax")
])

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 128)	100480
dense_1 (Dense)	(None, 10)	1290

=====

Total params: 101770 (397.54 KB)  
Trainable params: 101770 (397.54 KB)  
Non-trainable params: 0 (0.00 Byte)

```
In [5]: model.compile(optimizer="sgd",
    loss = "sparse_categorical_crossentropy",
    metrics=['accuracy'])

history = model.fit(x_train,y_train,validation_data=(x_test,y_test), epochs=10)
```



```
Epoch 1/10
1875/1875 [=====] - 8s 4ms/step - loss: 0.6570 - accuracy: 0.8362 - val_loss: 0.3618 - val_accuracy: 0.9010
Epoch 2/10
1875/1875 [=====] - 8s 4ms/step - loss: 0.3381 - accuracy: 0.9060 - val_loss: 0.2978 - val_accuracy: 0.9183
Epoch 3/10
1875/1875 [=====] - 7s 3ms/step - loss: 0.2908 - accuracy: 0.9184 - val_loss: 0.2647 - val_accuracy: 0.9267
Epoch 4/10
1875/1875 [=====] - 7s 4ms/step - loss: 0.2607 - accuracy: 0.9269 - val_loss: 0.2404 - val_accuracy: 0.9349
Epoch 5/10
1875/1875 [=====] - 7s 4ms/step - loss: 0.2378 - accuracy: 0.9341 - val_loss: 0.2219 - val_accuracy: 0.9372
Epoch 6/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.2189 - accuracy: 0.9392 - val_loss: 0.2066 - val_accuracy: 0.9413
Epoch 7/10
1875/1875 [=====] - 7s 4ms/step - loss: 0.2031 - accuracy: 0.9436 - val_loss: 0.1925 - val_accuracy: 0.9453
Epoch 8/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.1898 - accuracy: 0.9471 - val_loss: 0.1808 - val_accuracy: 0.9479
Epoch 9/10
1875/1875 [=====] - 7s 4ms/step - loss: 0.1781 - accuracy: 0.9504 - val_loss: 0.1723 - val_accuracy: 0.9514
Epoch 10/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.1679 - accuracy: 0.9532 - val_loss: 0.1637 - val_accuracy: 0.9531
```

```
In [6]: test_loss, test_acc=model.evaluate(x_test,y_test)
print("Loss = %.3f" %test_loss)
print("Accuracy = %.3f" %test_acc)


n = random.randint(0,9999)
plt.imshow(x_test[n])
plt.show()
predicted_value=model.predict(x_test)
plt.imshow(x_test[n])
plt.show()

print("Predicted value:", predicted_value[n])
```

313/313 [=====] - 1s 2ms/step - loss: 0.1637 - accuracy: 0.9531  
 Loss = 0.164  
 Accuracy = 0.953



313/313 [=====] - 1s 2ms/step



Predicted value: [1.4778496e-05 1.7555138e-05 1.0789034e-03 5.0535538e-05 1.3245249e-05  
 2.9411001e-04 1.6439888e-04 8.3090557e-09 9.9835998e-01 6.3908342e-06]

```
In [7]: #plotting the training accuracy
plt.plot(history.history["accuracy"])
plt.plot(history.history["val_accuracy"])
plt.title("Model Accuracy")
plt.ylabel("accuracy")
plt.xlabel("epoch")
plt.legend(["Train", "Validation"], loc = "upper right")
plt.show()

#plotting the training loss
plt.plot(history.history["loss"])
plt.plot(history.history["val_loss"])
plt.title("Model Loss")
plt.ylabel("loss")
plt.xlabel("epoch")
plt.legend(["Train", "Validation"], loc = "upper left")
plt.show()
```

